



Quality of Service Aware Data Stream Processing for Highly Dynamic and Scalable Applications


Isam Mashhour Hasan Al Jawarneh,
PhD Candidate, Cycle XXXII

Department of Computer Science and Engineering - DISI
University of Bologna, Italy
isam.aljawarneh3@unibo.it

Supervisor:
Prof. **Rebecca Montanari**

PhD program coordinator:
Prof. **Davide Sangiorgi**

Esame finale anno 2020
2nd April 2020



Outline

- **Motivation and thesis statement**
- Thesis contribution: a novel architecture for QoS-aware distributed processing of big spatial data
 - Processing and storage of static data
 - ✓ QoS-aware spatial **batch processing** engine
 - ✓ QoS-aware scalable **storage** for spatial data
 - Online Spatial Approximate Query Processing
 - ✓ QoS-aware **stream processing** engine for spatial data
 - ✓ Adaptive spatial **stream-static join** processor
- Conclusions and future works

Motivating Application Scenario


A mixed-workload scenario requiring at least:

- **Traffic Light Controller.** Actuator decides to change lights consistently for ambulance to pass
- **Smart Real-time Pathfinder.** Interactive navigation map for ambulances and other vehicles
- **Real-time Community Detector.** Identify volunteers' communities in the surroundings of the patient

➤ Primitive geospatial queries

- Proximity queries
- Spatial join
- Spatial clustering
- Spatial geo-statistics.
- k -Nearest Neighborhoods)


participatory healthcare




Geospatial data representation

The problem of multidimensionality

- Big geospatial data management is challenging
- A spatial point is parametrized and represented as coordinates (longitude and latitude)
- GPS is rarely 100% accurate, susceptible to acceptable error-bounds
- **Geometry** inherent in the data will be **lost** by such a transformation
- Spatial reconstruction is **expensive**



parameterizing

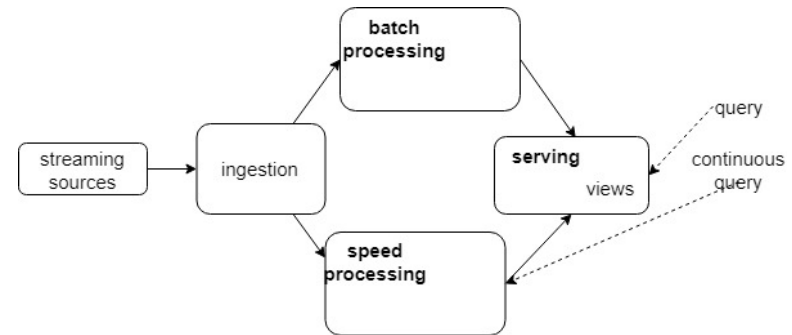


Longitude	Latitude
11.3709	44.5185
11.4081	44.4963
11.3477	44.499

- Dimensionality reduction. **Geohash** is a geospatial encoding
- Generates a single-dimensional representation as a string that encompasses a geographical meaning

Related works

- Few ad-hoc fixes, patches and glues for spatial data management:
 - Spatial partitioning methods for **batch** processing
 - Spatial query **optimizers**
 - Spatial **stream** processing
- They do not collectively form a comprehensive architecture for mix workload demanding dynamic applications
- Drawbacks of **Lambda** architecture
 - Not attuned to the spatial characteristics of data
 - QoS-awareness is not incorporated transparently
 - Users need to explicitly reason about underlying logistics



Lambda [1]

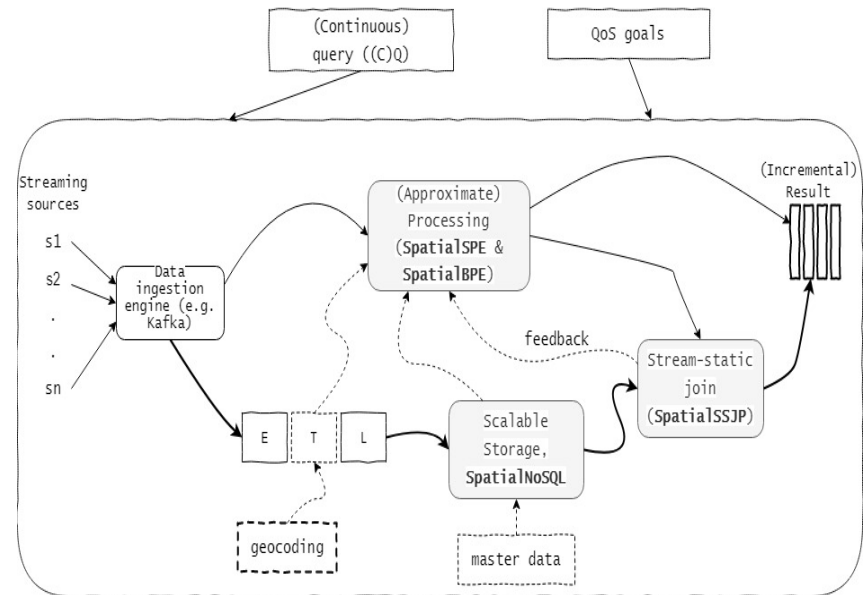
[1] N. Marz and J. Warren, *Big Data: Principles and Best Practices of Scalable Real-Time Data Systems*. New York; Manning Publications Co., 2015.

SpatialDSMS: a platform for geo big data

- Design guidelines


- A platform operating atop best-in-breed Cloud computing representatives
- Same multidimensionality **reduction** approach for all workloads: storage, batch and stream processing
- All layers collaborate **synergistically** in serving a mixed workload dynamic application scenario with QoS guarantees
- Design goals include **modularity** and **composability**
- Our new layer becomes the new codebase for QoS awareness

Spatial data stream management system



- QoS goals include

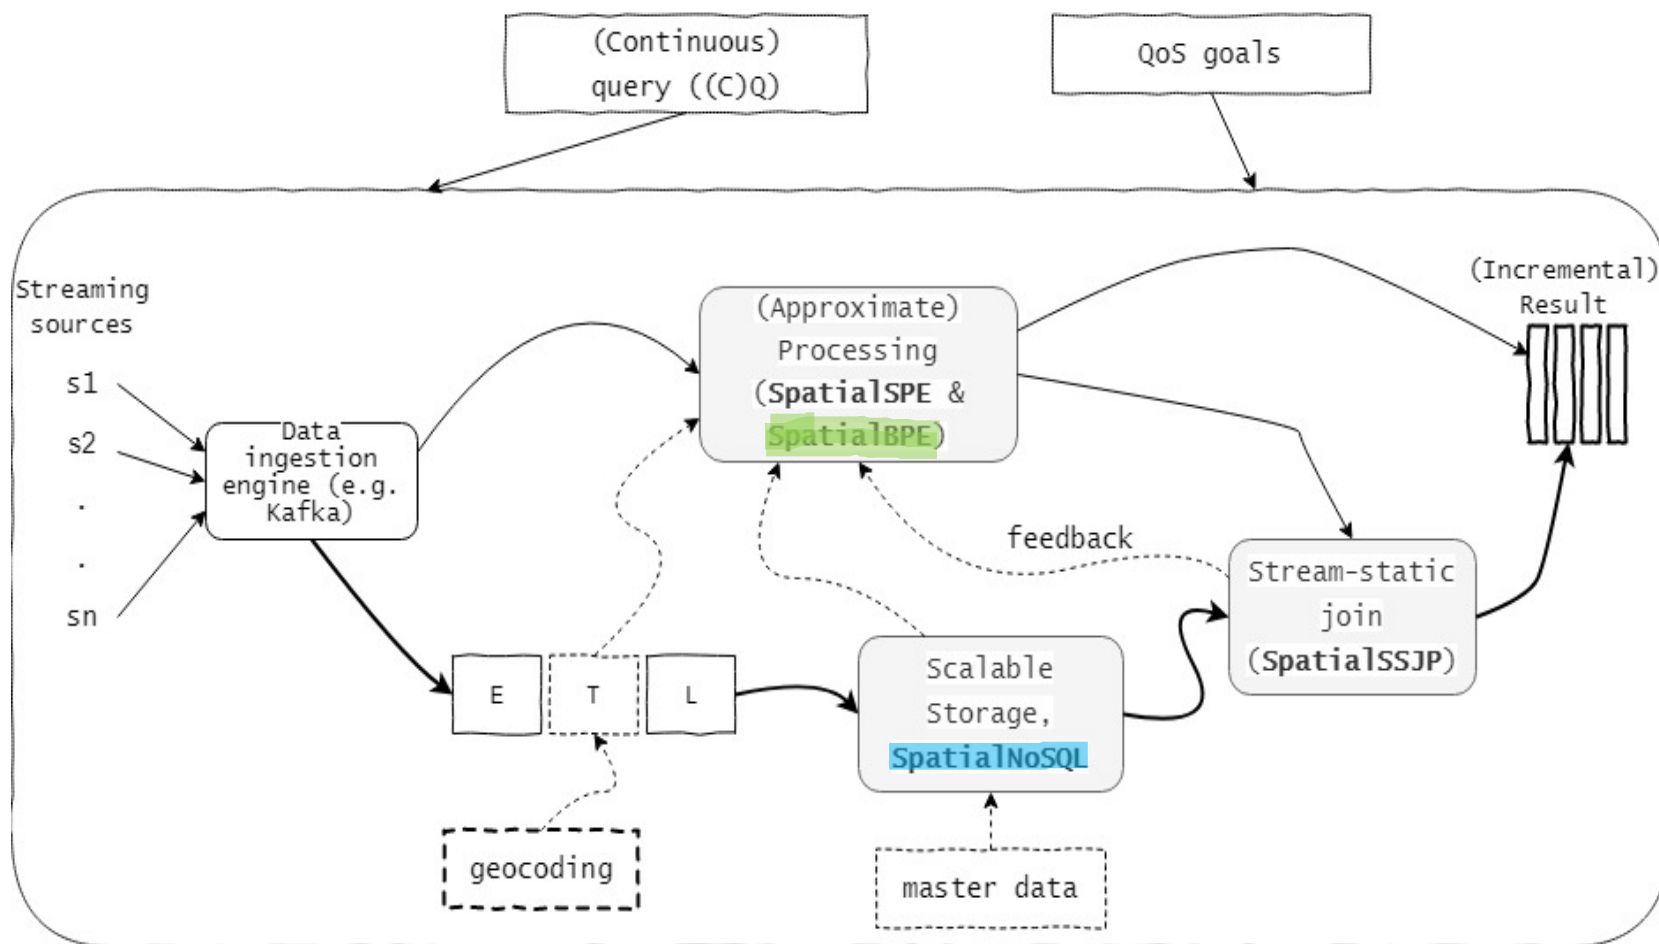
- ✓ Low latency
- ✓ High throughput
- ✓ Maximum resource utilization
- ✓ Maximum accuracy



Processing and storage of static data: optimization pillars

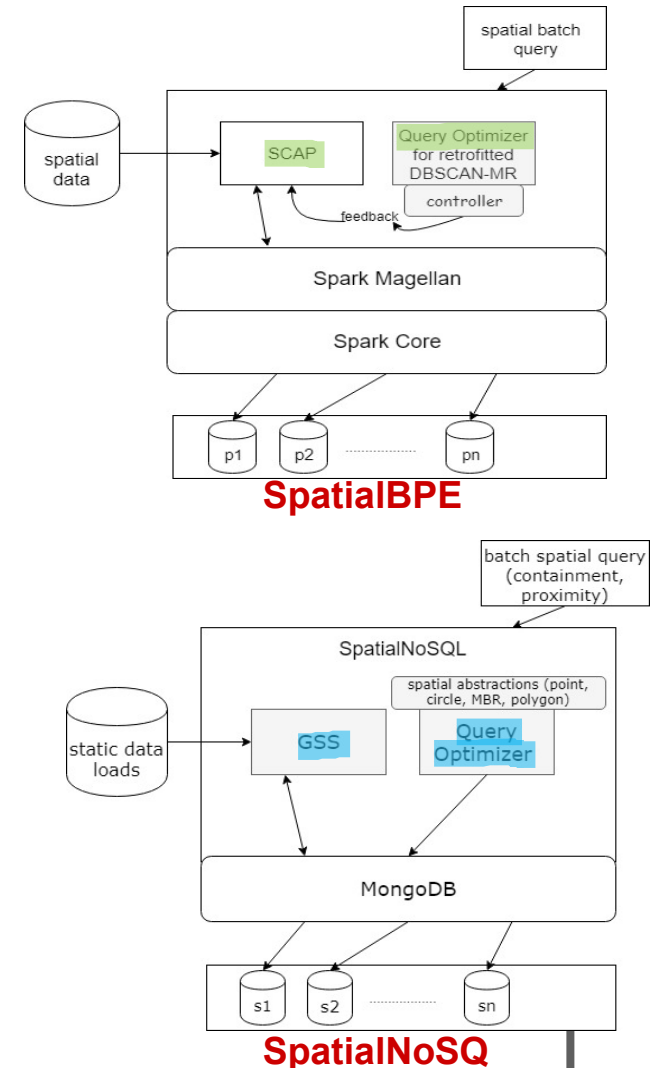
- Two aspects are important for optimizing the processing and storage of spatial data-at-rest:
 - Spatial **data partitioning**
 - ✓ Splitting data to parallelly internetworked processing (or storage) worker nodes
 - ✓ Sharding in NoSQL scalable storage is analogous
 - Spatial **query optimizers**
 - ✓ Optimizing spatial query by selecting the best performing query plan
- Spatial data partitioning goals
 - Load balancing
 - Spatial Data Locality (SDL) preservation
 - Boundary Spatial Objects (BSO) minimization

QoS-aware static big spatial data processing



SpatialBPE and SpatialNoSQL: Overview

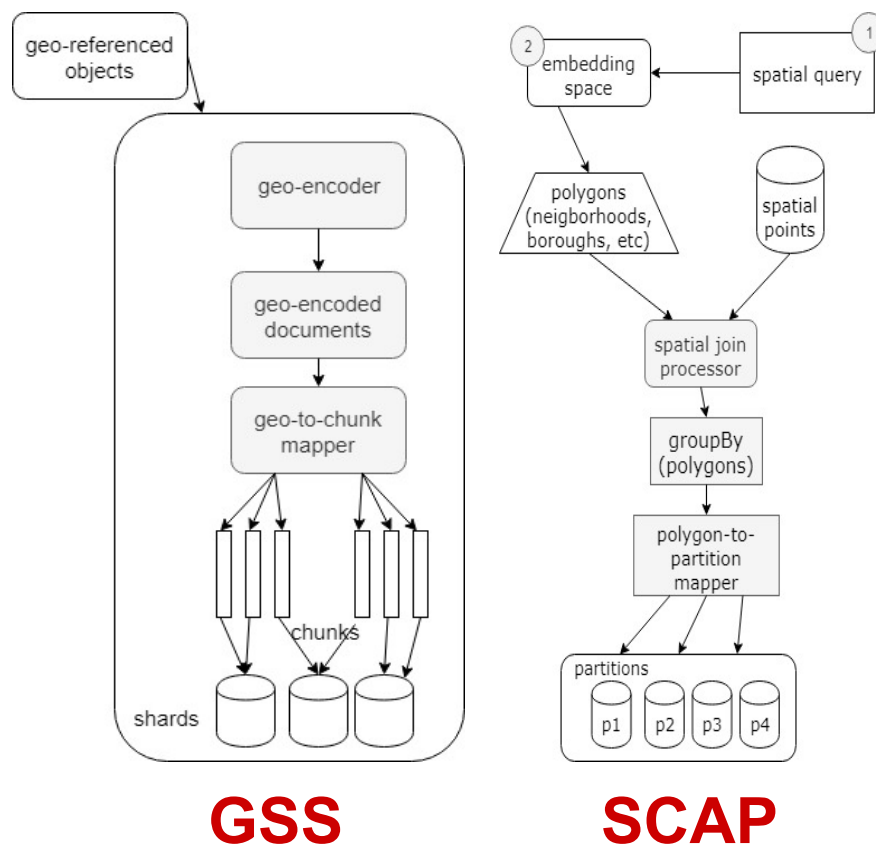
- Spatial QoS-aware batch processing system
- Two component:
 - Custom **partitioning method**: SCAP for batch systems and GSS for NoSQL storage
 - Spatial **query optimizer**, exploiting SCAP and GSS for clustering, spatial join, proximity and containment queries



Spatial Co-locality-aware Partitioner (SCAP) & Geospatial Sharding Scheme (GSS)

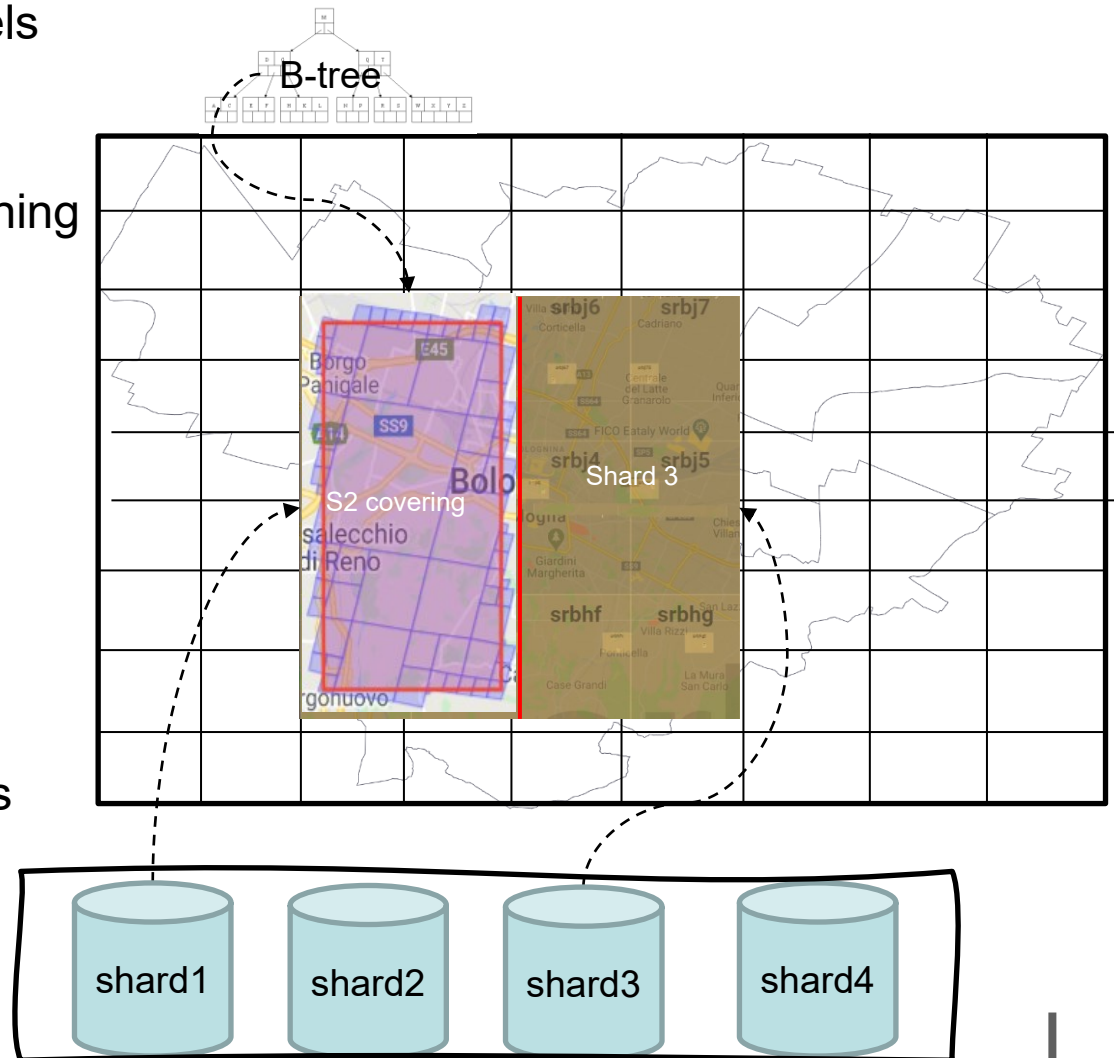
- **SDL** preservation is a priority
- **BSOs** and **load balancing** to a lesser extent

- **Clump** geometrically co-located objects into single chunks
- **Split** overloaded chunks
- **Map** chunks to partitions




Spatial query optimizer for NoSQL: Overview

- We have created a two-level indexing scheme
- MongoDB router forwards requests to few shards, pruning the search space
- **Overlay** the embedding space with a fixed-grid network
- **Generate** a geohash covering and a list of interacting points
- **Generate** S2 sub-coverings and retrieve interacting points
- **Impose** B-tree index on sub-coverings




Results: achieving QoS goals for clustering



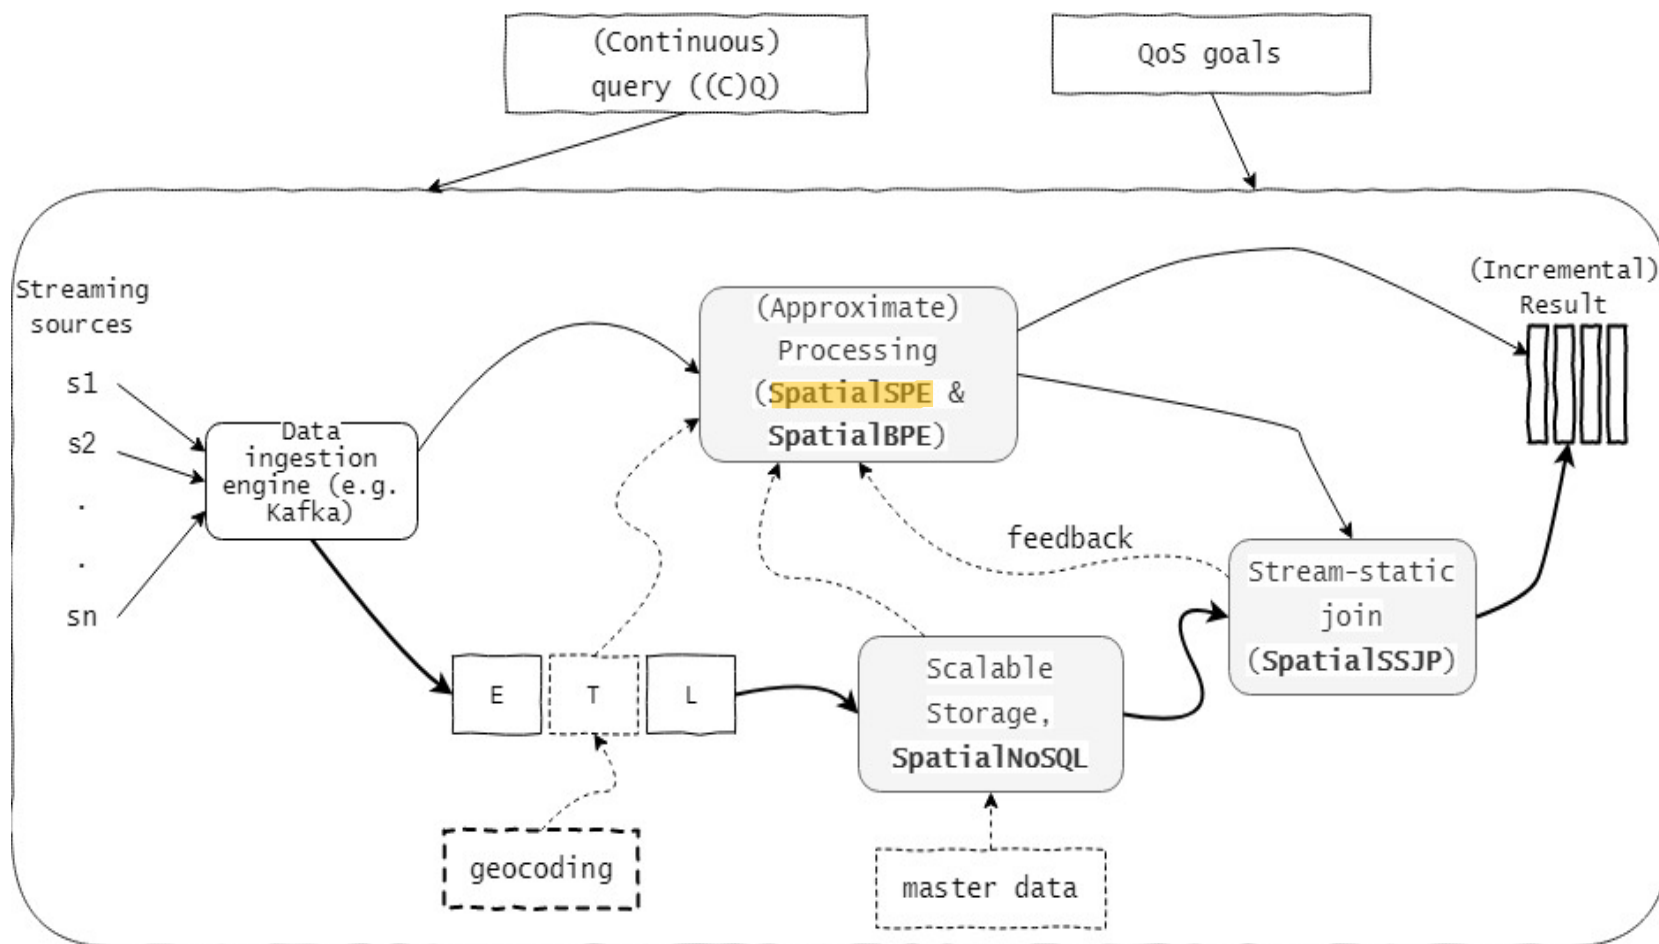
- **NY City** taxicab itinerary datasets a cohort of approximately 150k
- 150k spatial data points that were collected through the **ParticipAct** project by our group in Bologna, Italy
- Tweak geohash from 30 to 35
- Total running time has been decreased

Results: containment test




- **NY City** taxicab trips datasets, a cohort of two months (around three million units)
- Geohash 30, our geohash-based query optimizer searches **three shards only**
- Documents and keys examined using our optimizer are **less** than those using plain MongoDB

QoS-aware interactive big spatial data processing



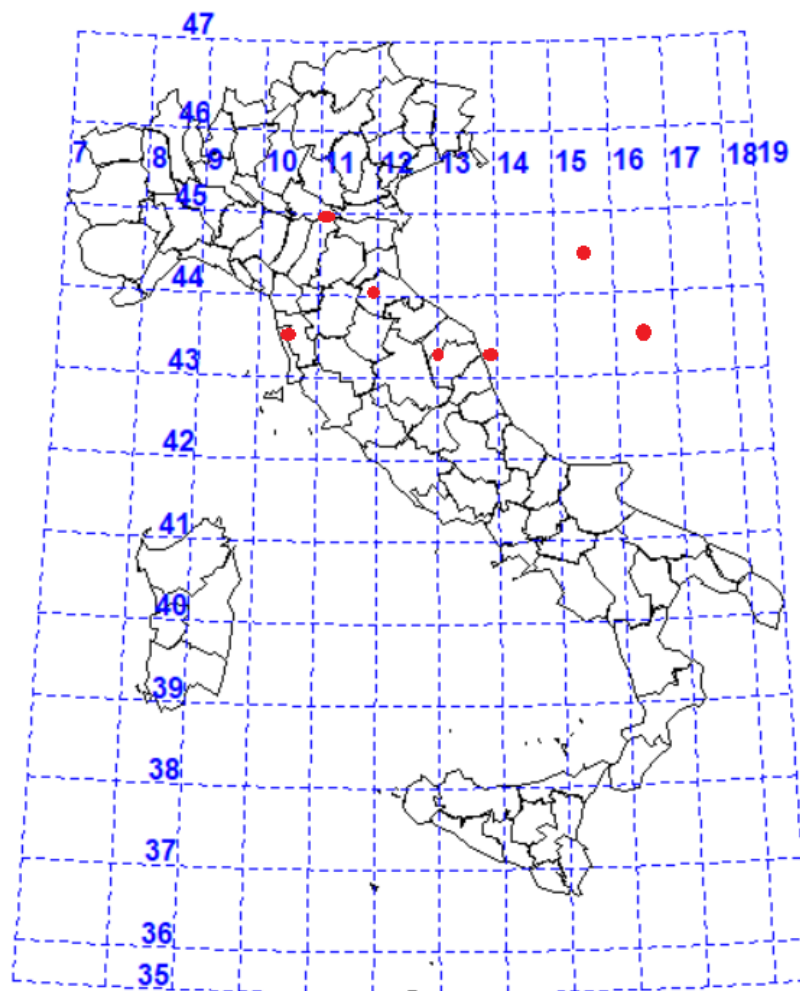
Parametrized spatial data



Embedding area polygons

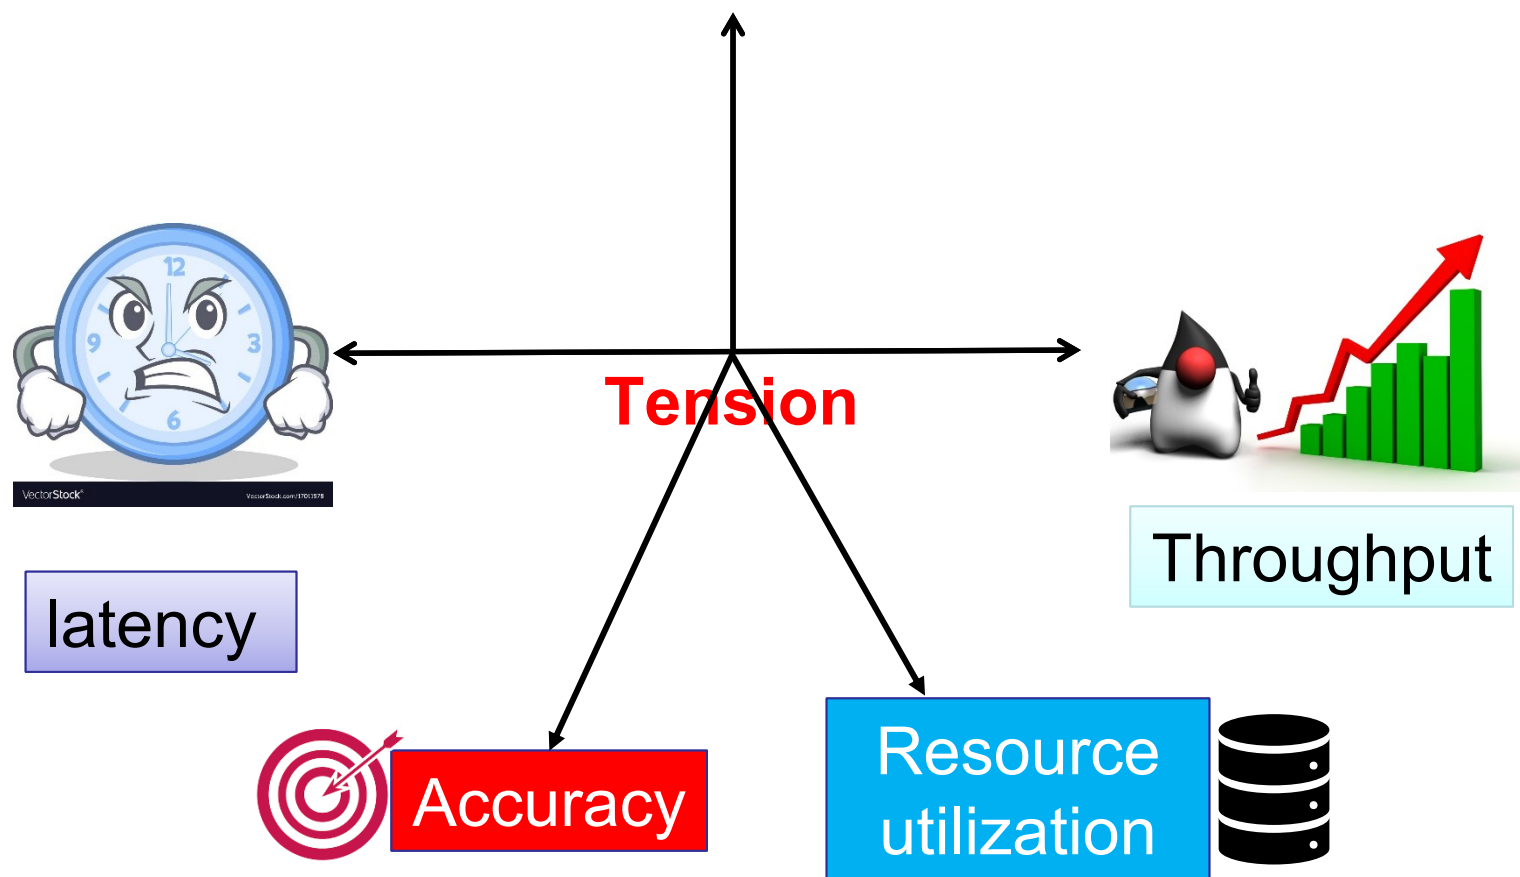


Overlaying maps




QoS Tension


Spatial (**Approximate**) Query Processing (S(**A**)QP)



Spatial Approximate Query Processing (SAQP)

- Stream Processing Engines (SPEs) are confronted with complex challenges
 - Fast arriving streaming workloads
 - Temporal arrival rate fluctuation and skewness
- Can we do better?
 - After 1 second, we obtain a 99.5 accurate early result, which is satisfactory for decision making, which then makes the final exact result not needed

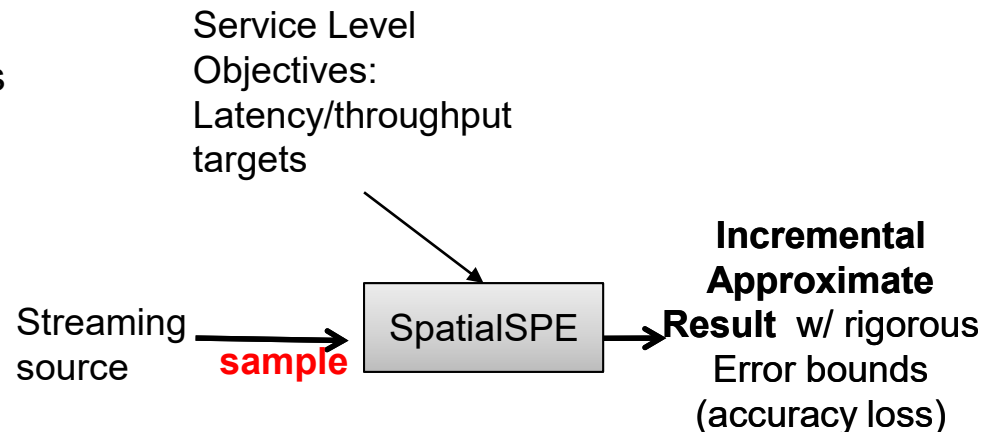




SpatialSPE

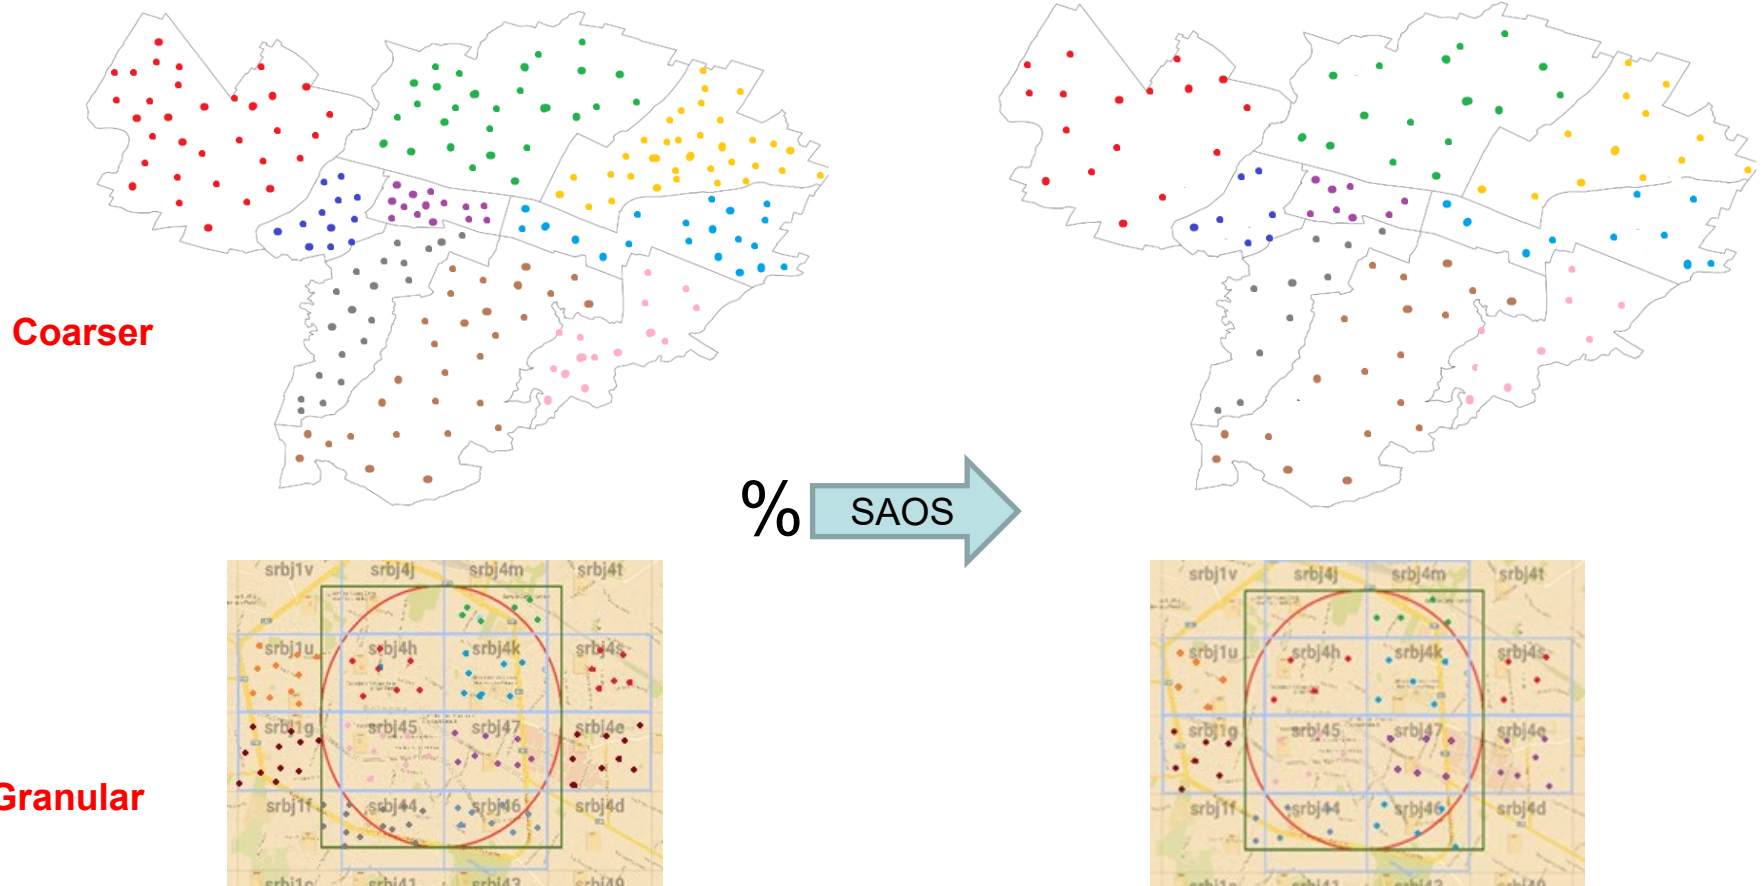
- Spatial data maintain spatial trends that affect the observed responses
 - **spatially representative samples**
→ selecting spatially well-spread out samples positively affects the accuracy of estimators (average, median, etc.).
 - A sample constitutes a scaled-down ('microcosm') version, mirroring characteristics of the population it is representing
- **Example Continuous Query (CQ).**
“measuring the average trip distance travelled by taxis from each borough in NYC, United States”
- Sampling fractions are the same for all constituent stratum
- CQ is **incrementalized**.

Computing over a sample instead of the whole population



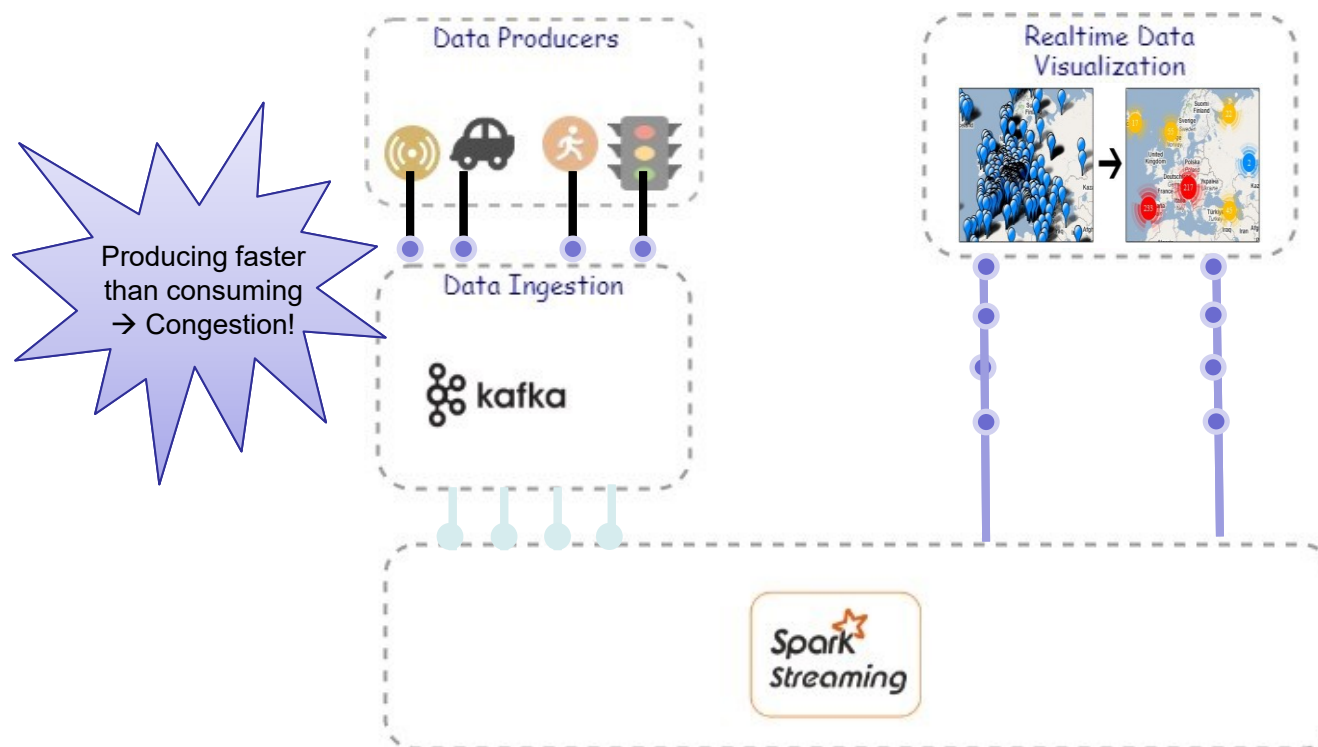
SpatialSPE overview

Spatial Aware Online Sampling (SAOS): overview

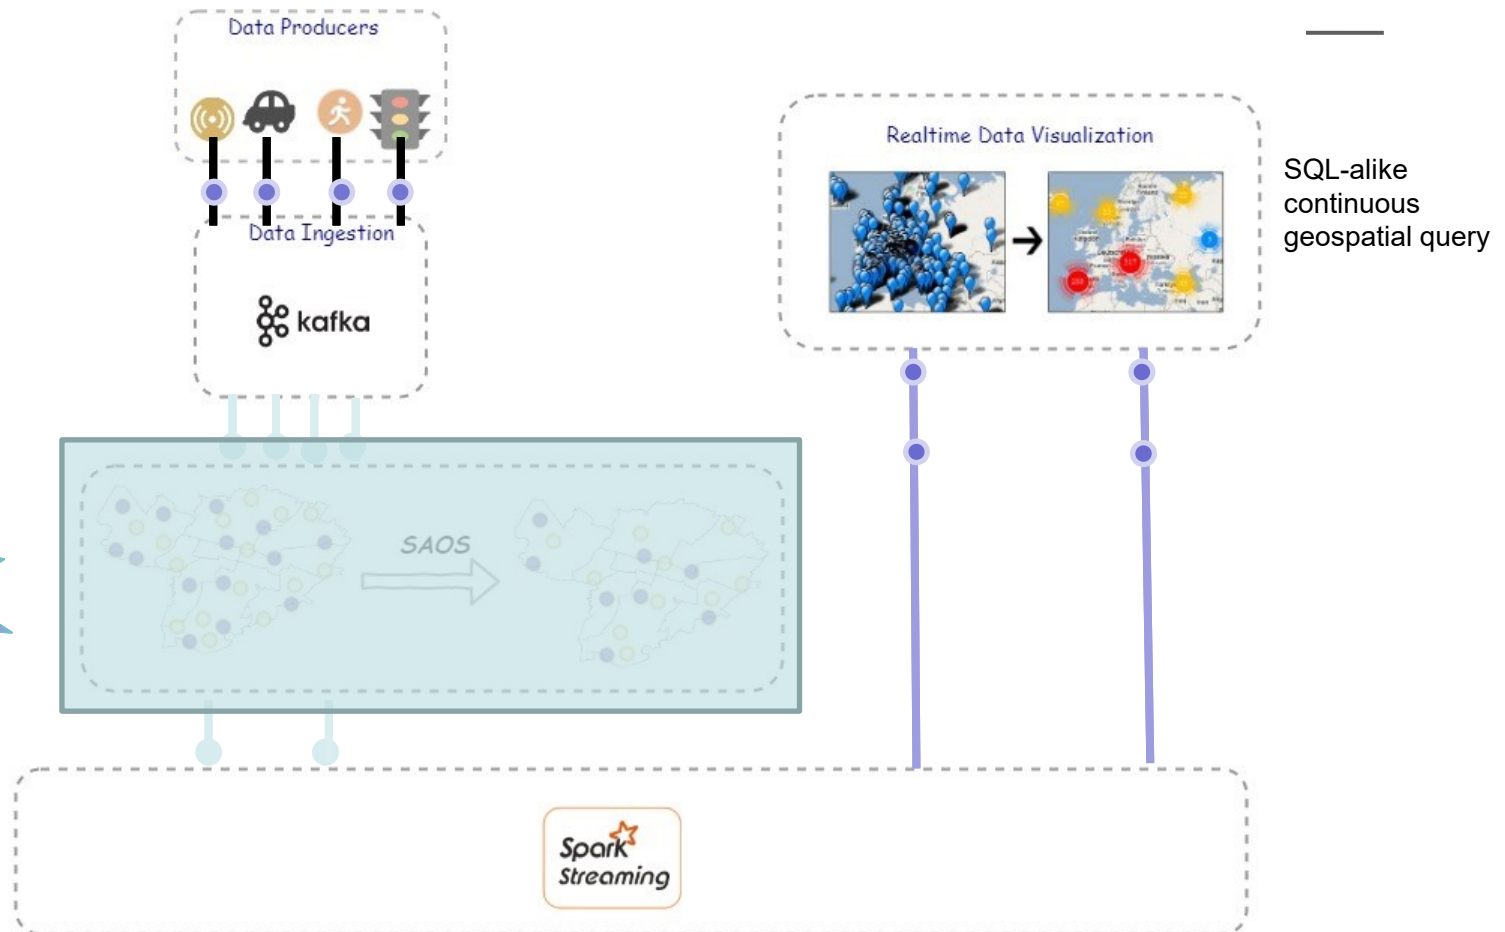


- Nearby points share the same geohash prefixes
- **SAOS** focuses on **SDL preservation**


Typical pipeline architecture w/o SAOS



The improved architecture w/ SAOS

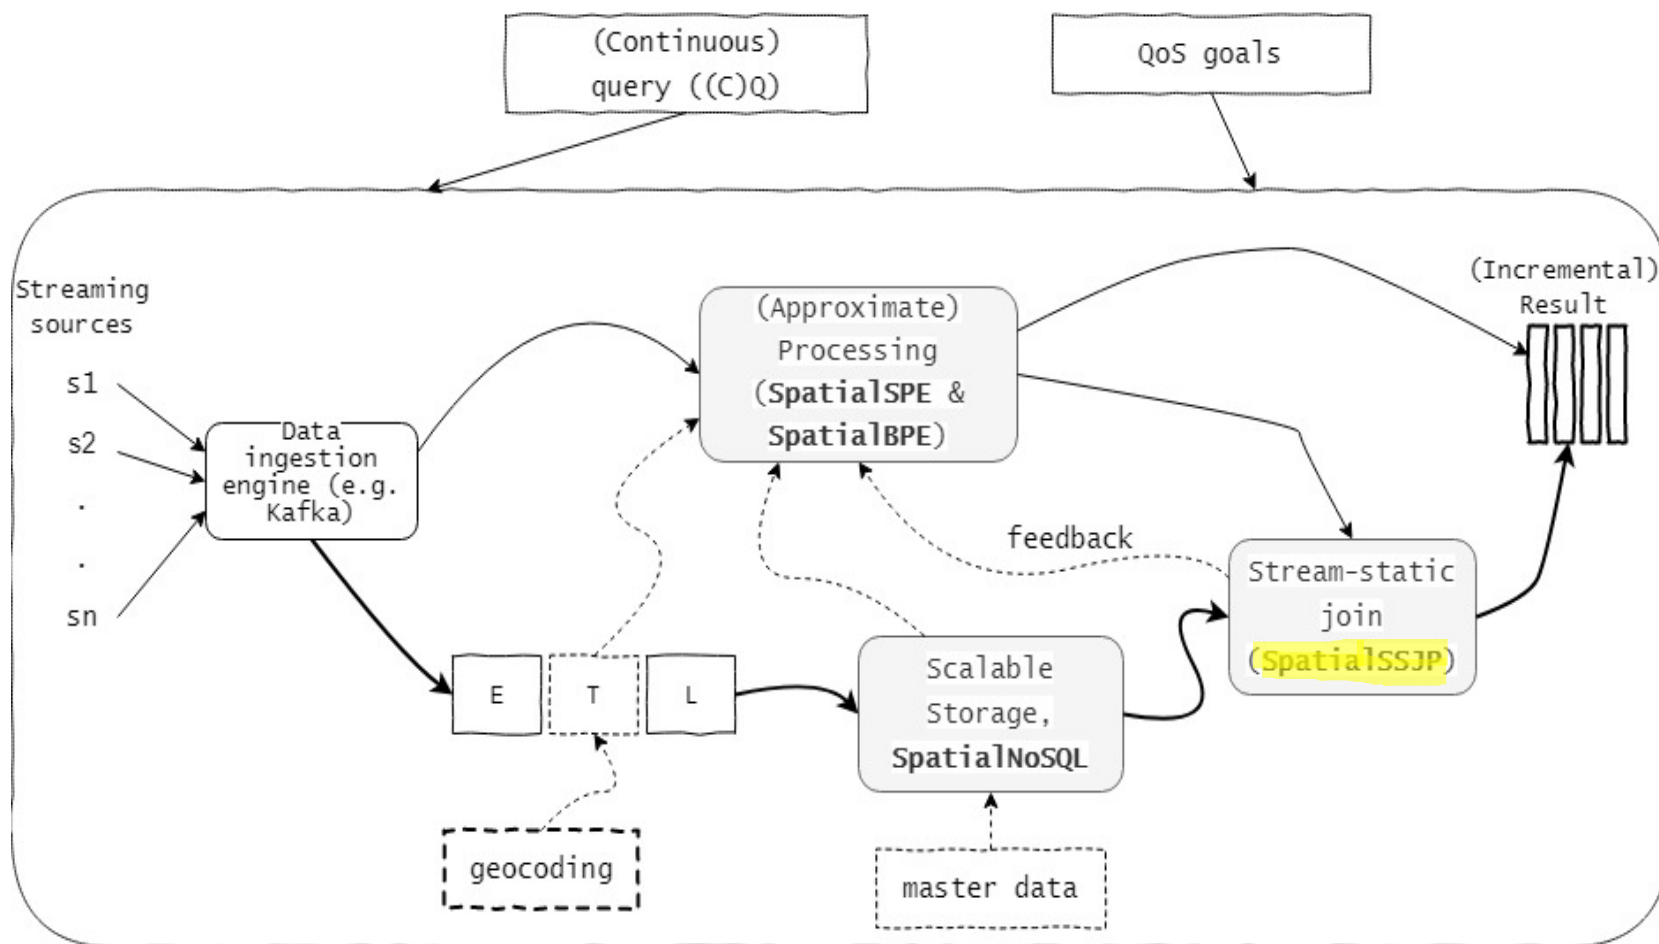


Results: achieving accuracy targets



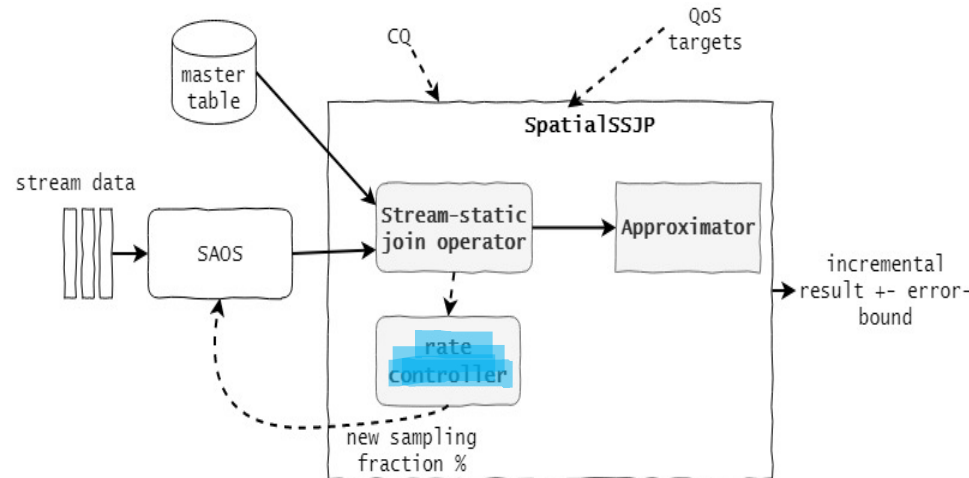
- **NY City** green taxicab trips datasets , where we select a big cohort representing six months (almost nine million tuples)
- SAOS outperforms SpSS-based SRS for all precision settings (30 and 35), for both measures, accuracy loss and relative error.
- SAOS have bigger accuracy loss for geohash precision 35 , compared to SAOS accuracy loss at geohash precision 30.

SpatialDSMS: a platform for geo big data



SpatialSSJP Overview

- Adaptive QoS- and Spatial-aware framework for processing spatial stream-static joins
 - **Continuous spatial query** requiring stream-static join and a **query budget** are served
 - **Rate controller** computes sampling fraction
 - **SAOS** samples based on the fraction
 - Stream-static **join operator** result is fed to an approximator
 - **Approximator** serves an incremental result with rigorous error-bounds
- Hybridizing a novel rate controller with SAOS from SpatialSPE
- Either latency or accuracy QoS guarantees



```

SELECT point p, polygon po, avg(tripDistance)
FROM Stream S, MasterTable M
WHERE S.key = M.key AND (p WITHIN po)

GroupBy neighborhood

Latency 120 MS

OR

e 0.03 CL 95%

```

Rate controller

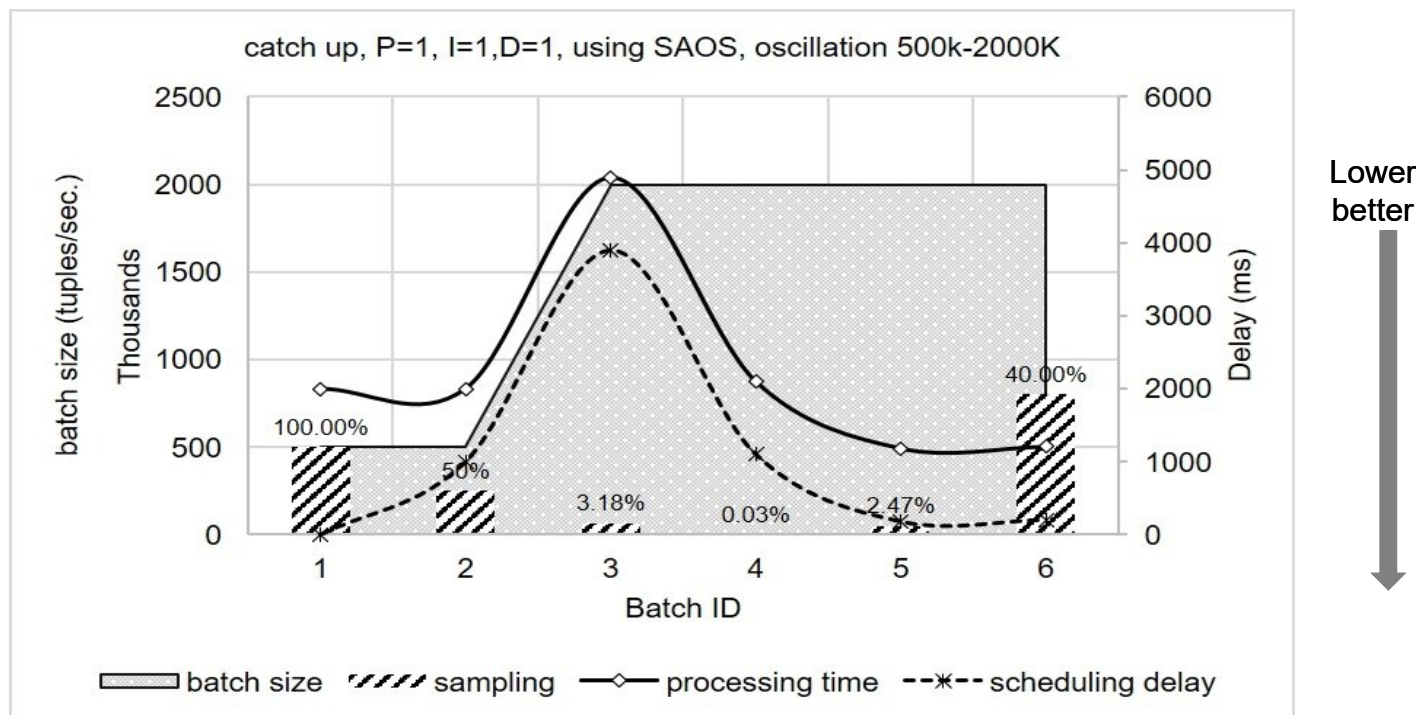
Latency-aware rate controller

- PID controller is a control loop feedback mechanism that calculates an error value
- Three terms (PID):
 - **Proportional**: present error
 - **Integral**: historical errors
 - **Derivative**: future errors
- After each trigger, the new rate is calculated
 - $rate_{new} = rate_{latest} - ((P.err) + (I.err_{hist}) + (D.err_d)) = SP - PV$

Accuracy-aware rate controller

- 'margin of error' specified as a QoS target
- For SAOS, we depend on theory of stratification, for 95% confidence level, we compute the sampling fraction:
 - $n = 3.84 * (v/e_{des}^2)$

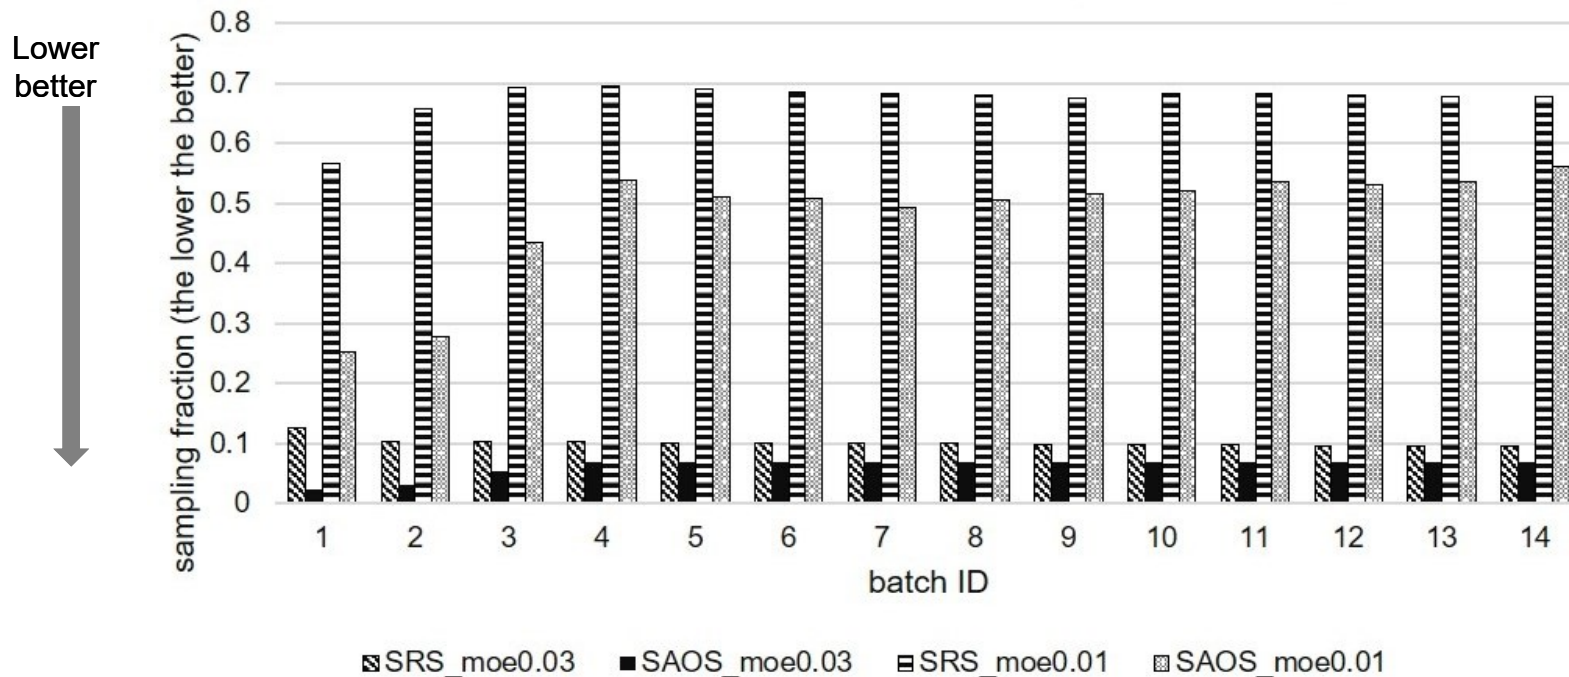
Results: achieving latency goals




- **NY City** green taxicab trips datasets , where we select a big cohort representing six months (almost nine million tuples)
- SpatialSSJP could survive brutal spikes in the data stream arrival rates

Results: achieving accuracy goals

gain of using SAOS against SRS in SpatSSJP, margin of error 0.03 and 0.01



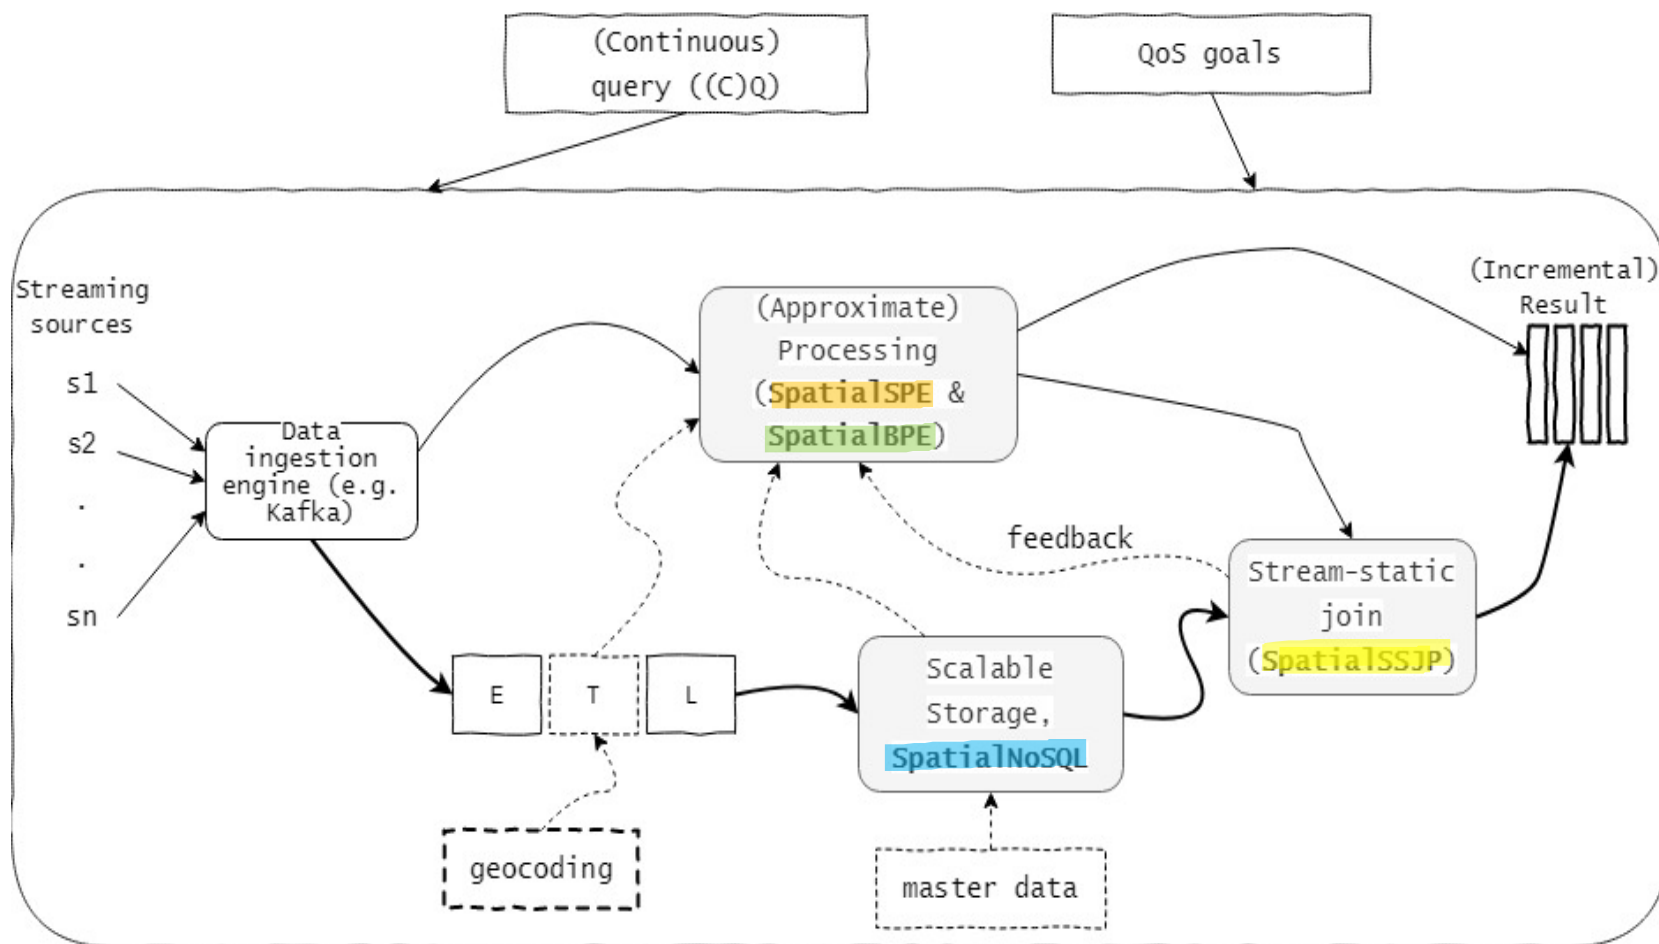
- SRS-based or SAOS achieve **accuracy** targets
- However, SAOS requires, on average, **less sampling fractions** → lower latency and higher resource utilization
- **Restrictive** cases (margin of error 0.01 imply more sampling fractions)




Outline

- Motivation and thesis statement
- Thesis contribution: a novel architecture for QoS-aware distributed processing of big spatial data
 - Processing and storage of static data
 - ✓ QoS-aware spatial batch processing engine
 - ✓ QoS-aware scalable storage for spatial data
 - Online Spatial Approximate Query Processing
 - ✓ QoS-aware stream processing engine for spatial data
 - ✓ Adaptive spatial stream-static join processor
- Conclusions and future works


SpatialDSMS: a platform for geo big data






Summary of contributions

- Avalanches of geospatial data present businesses with formidable challenges
- “one-size-fits-all” does not hold true, data-at-rest need to be **combined** with data-in-motion for deeper insights
- Current systems do not **natively** offer QoS awareness as a transparent underlying layer for processing streams of geo-referenced data
- Constituent parts should **operate synergistically** in achieving QoS goals
- We have designed a QoS Aware DSMS for geo-referenced huge amounts of streaming data loads in **highly dynamic scenarios**, SpatialDSMS:
 - ✓ QoS-aware spatial **batch processing** engine
 - ✓ QoS-aware **scalable storage** for spatial data
 - ✓ QoS-aware **stream processing** engine for spatial data
 - ✓ Adaptive spatial **stream-static join** processor
- A modular architecture that streamlines the orchestration between the constituent sub-systems




Applicability in diverse domains

- QoS -aware optimizations we have provided in this thesis are in no way exhaustive
- They constitute precursors for other domain-specific optimizations
- mixed workloads that are easily composable by mixing some of the services we provide
 - **Real-time traffic control**
 - We offer **baselines** for building a fully-functional real-time traffic control system: *flow rates*, *occupancy* and *density*
 - We support **incrementalization** for a primitive set of spatial statistics
 - **Spatial online stream clustering**
 - Composability : using the baseline primitives we support
 - Online clustering algorithms work by combining two phases:
 - **Online**: *incrementally* cluster data points based on proximity, forming **micro-clusters**
 - **Offline**: **macro-clustering**, forming actual clusters in batch mode using an advanced clustering algorithm




Open research areas

- **Designing** online spatial-aware data partitioning schemes
 - Taking spatial partitioning goals online imposes challenges that do not affect batch partitioning
- **Offloading** sequential jobs to **Fog** nodes
 - Sending endlessly huge amounts of geo-referenced loads to the **cloud** which could be detrimental in low-latency applications
 - Example, an online **sampler** can be pushed upstream near the Edge, achieving better latency QoS goals




Relevant publications

- [1] **I. M. Al Jawarneh**, P. Bellavista, L. Foschini and R. Montanari, "Spatial-aware approximate big data stream processing," in 2019 IEEE Global Communications Conference (GLOBECOM), 2019, pp. 1-6.
- [2] **I. M. Al Jawarneh**, P. Bellavista, A. Corradi, L. Foschini, R. Montanari and A. Zanotti, "In-memory spatial-aware framework for processing proximity-alike queries in big spatial data," in 2018 IEEE 23rd International Workshop on Computer Aided Modeling and Design of Communication Links and Networks (CAMAD), 2018, pp. 1-6.
- [3] **I. M. Al Jawarneh**, P. Bellavista, F. Casimiro, A. Corradi and L. Foschini, "Cost-effective strategies for provisioning NoSQL storage services in support for industry 4.0," in 2018 IEEE Symposium on Computers and Communications (ISCC), 2018, pp. 1227.
- [4] **I. M. Aljawarneh**, P. Bellavista, C. R. De Rolt and L. Foschini, "Dynamic identification of participatory mobile health communities," in Cloud Infrastructures, Services, and IoT Systems for Smart Cities. Springer, 2017, pp. 208-217.
- [5] **I. M. Aljawarneh**, P. Bellavista, A. Corradi, R. Montanari, L. Foschini and A. Zanotti, "Efficient spark-based framework for big geospatial data query processing and analysis," in 2017 IEEE Symposium on Computers and Communications (ISCC), 2017, pp. 851-856.
- [6] **I. M. Aljawarneh**, P. Bellavista, A. Corradi, L. Foschini, R. Montanari, "Efficient QoS-Aware Spatial Join Processing for NoSQL Scalable Storage Frameworks". 2020. **Submitted**.
- [7] **I. M. Aljawarneh**, P. Bellavista, A. Corradi, L. Foschini, R. Montanari, ". SpatialSSJP: QoS-Aware Adaptive Approximate Stream-Static Spatial Join Processor". 2020. **Submitted**.
- [8] **I. M. Aljawarneh**, P. Bellavista, A. Corradi, L. Foschini, R. Montanari, " Locality-Preserving Spatial Partitioning Scheme for Quality Spatial Analytics in Distributed Main Memory Frameworks".2020. **Submitted**.
- [9] **I. M. Aljawarneh**, P. Bellavista, A. Corradi, L. Foschini, R. Montanari. "QoS-Aware Optimizations for Big Geospatial Data Management - A Survey". 2020. **Submitted**.



Other publications


- [10] **I. M. Al Jawarneh**, P. Bellavista, A. Corradi, L. Foschini, R. Montanari, J. Berrocal and J. M. Murillo, "A Pre-Filtering Approach for Incorporating Contextual Information Into Deep Learning Based Recommender Systems," IEEE Access, vol. 8, pp. 40485-40498, 2020.
- [11] S. Bertacchi, **I. M. Al Jawarneh**, F. I. Apollonio, G. Bertacchi, M. Cancilla, L. Foschini, C. Grana, G. Martuscelli and R. Montanari, "SACHER project: A cloud platform and integrated services for cultural heritage and for restoration," in Proceedings of the 4th EAI International Conference on Smart Objects and Technologies for Social Good, 2018, pp. 283-288.
- [12] **I. M. Al Jawarneh**, P. Bellavista, L. Foschini, R. Montanari, J. Berrocal and J. M. Murillo, "Toward privacy-aware healthcare data fusion systems," in International Workshop on Gerontechnology, 2018, pp. 26-37.
- [13] **I. M. Al Jawarneh**, P. Bellavista, F. Bosi, L. Foschini, G. Martuscelli, R. Montanari and A. Palopoli, "Container orchestration engines: A thorough functional and performance comparison," in ICC 2019-2019 IEEE International Conference on Communications (ICC), 2019, pp. 1-6.
- [14] **I. M. Al Jawarneh**, P. Bellavista, L. Foschini, G. Martuscelli, R. Montanari, A. Palopoli and F. Bosi, "Qos and performance metrics for container-based virtualization in cloud environments," in Proceedings of the 20th International Conference on Distributed Computing and Networking, 2019, pp. 178-182.
- [15] P. Bellavista, J. Berrocal, A. Corradi, S. K. Das, L. Foschini, **I. M. Al Jawarneh** and A. Zanni, "How Fog Computing Can Support Latency/Reliability-Sensitive IoT Applications: An overview and a Taxonomy of State-Of-The-Art Solutions," 2019.



A Tutti Voi... Grazie!

to the real
heroes!





Quality of Service Aware Data Stream Processing for Highly Dynamic and Scalable Applications

A Tutti Voi... Grazie!

By:

Isam Mashhour Hasan Al Jawarneh,

PhD Candidate, Cycle XXXII

Supervisor:
Prof. **Rebecca Montanari**

Department of Computer Science and Engineering - DISI
University of Bologna, Italy
isam.aljawarneh3@unibo.it

PhD program coordinator:
Prof. **Davide Sangiorgi**

Esame finale anno 2020

2nd April 2020